

COTS-Based Fault Tolerance in Deep Space: Qualitative and Quantitative Analyses of a Bus Network Architecture*

Ann T. Tai
IA Tech, Inc.
10501 Kinnard Avenue
Los Angeles, CA 90024

Savio N. Chau Leon Alkalai
Jet Propulsion Laboratory
California Institute of Technology
Pasadena, CA 91109

Abstract

Among the COTS applications in the X2000 architecture for deep-space missions, the use of commercial bus standards is the highest-payoff COTS application since a bus interface has a global impact and enabling effect on system cost and capability, respectively. While COTS bus standards enable significant cost reductions, it is a great challenge for us to deliver a highly-reliable long-term survivable system employing COTS standards that are not developed for mission-critical applications. The spirit of our solution to the problem is to exploit the pertinent standard features of a COTS product to circumvent its shortcomings, though these standard features may not be originally designed for highly reliable systems. In this paper, we discuss our experiences and findings on the design and assessment of an IEEE 1394 compliant fault-tolerant bus architecture. We first derive and qualitatively analyze a “stack-tree topology” that not only complies with IEEE 1394 but also enables the implementation of a fault-tolerant bus architecture without node redundancy. We then present a quantitative evaluation that demonstrates significant reliability improvement from the COTS-based fault tolerance.

1 Introduction

The X2000 system architecture that has been developed by NASA/JPL is a distributed, scalable, fault-tolerant avionics architecture for multiple deep-space missions [1, 2, 3]. The architecture is currently the baseline for the Europa Orbiter mission which is scheduled for launch in year 2003 [4]. In the X2000 architecture, the multiple computing nodes and devices are symmetric, meaning that the roles of computing nodes are interchangeable while devices are treated as intelligent nodes in the network. Moreover, driven by NASA’s *faster, better, cheaper* space mission philosophy, this architecture stresses on using *commercially-off-the-shelf* (COTS) products, standards and intellectual properties (IPs) to reduce development and recurring costs.

Among the COTS applications in the X2000 architecture, the use of commercial bus standards is the highest-payoff COTS application. This is due to the fact that, unlike a system component or subsystem, a bus interface is normally used throughout the entire system architecture and thus has a *global* impact and enabling effect on system cost and capability, respectively. After extensive

survey and rigid evaluation efforts, the Peripheral Component Interface (PCI) [5], IEEE 1394 [6, 7], and I2C [8] have been selected to implement the local computer bus, system bus, and engineering bus, respectively, in the X2000 architecture [9]. While COTS bus standards enable significant cost reductions, it is a great challenge for us to deliver a highly-reliable long-term survivable system employing COTS standards that are not developed for mission-critical applications. The central purpose of this paper is to report our experiences and findings on the design and assessment of the IEEE 1394 compliant fault-tolerant bus network for the X2000 architecture.

The IEEE 1394 standard has two implementations, namely, *cable implementation* and *backplane implementation* [6]. The backplane implementation that adopts the multi-drop topology has been investigated by the aerospace industry [10]. Yet, we have selected the cable implementation for the X2000 architecture because of its higher data rate, lower power consumption and significantly more substantial commercial support. Nonetheless, in terms of fault tolerance, the tree topology criterion for the cable implementation makes it formidable for a non-redundant bus network to tolerate node/link failures. Namely, any single node or link failure will result in tree partitioning such that fault-tolerant routing will not be possible. A brute-force approach is to duplicate the circuitry of each node and to cross-strap them to the bus network. However, this approach is not allowed in the X2000 architecture because the significant increases in power consumption, mass, and volume will violate their constraints. Although IEEE 1394 has been enjoying fast-growing popularity, how to circumvent its limitations in fault tolerance for implementing highly reliable systems has not yet received enough attention. The spirit of our solution to the problem is to exploit the standard features of COTS products and standards in an innovative manner to circumvent their shortcomings, though these standard features may not be originally designed for fault tolerance. In this paper, we derive and qualitatively analyze a special realization of tree topology, called the “stack-tree topology,” which not only complies with the IEEE 1394 standard but also enables the implementation of a fault-tolerant bus architecture. In particular, the bus network architecture based on stack-tree topology enables us to exploit a standard feature of IEEE 1394 called “port-disable” [11], which ensures the normal communications among nodes to continue so long as no non-clustered multiple cut-type failures occur (see Section 3). We also present a reliability evaluation that quantitatively compares the bus network architectures based on the stack-tree topology and

*The research described in this paper was carried out by the Jet Propulsion Laboratory, California Institute of Technology, under a contract with the National Aeronautics and Space Administration.

its variants, namely, CST_S , CST_D and CST_R . Due to the complex failure scenarios of the network architecture CST_R , to take into account exhaustively and exclusively for all the failure scenarios based on straightforward use of combinatorics would be very difficult and inefficient. Therefore, we construct an analytic model based on recursive functions that traverse the entire network and enumerate all the concerned scenarios, yielding exact solution for the reliability measure. The evaluation results demonstrate that the proposed COTS-based fault-tolerant bus architecture leads to significant reliability gains for long-life deep-space missions.

The remainder of the paper is organized as follows. In Section 2, we explain why we have selected IEEE 1394, and describe its advantages and restrictions we exploit and circumvent, respectively. In Section 3, we elaborate and qualitatively analyze the stack-tree topology that complies with IEEE 1394 and facilitates the use of the “port-disable” feature for fault tolerance, followed by Section 4 which presents the quantitative methods and results of reliability evaluation for the resulting fault-tolerant bus architecture. In the concluding remark, we discuss the significance of this effort.

2 IEEE 1394: Its Selection, Advantages, and Restrictions

In the process of selecting the high-speed and low-power buses, we examined many commercial bus interfaces. The candidates included IEEE 1394, Fiber Channel, Universal Serial Bus (USB), Fast Ethernet, Serial Fiber Optic Data Bus (SFODB), ATM, Myrinet, FDDI, AS1773, and SPI. Many of them (e.g., USB, AS1773, and SPI) fail to meet the requirements on data rate; some of them are not suitable for real-time applications because of the indeterminacy of bus latency; whereas others have high power consumption which is unacceptable for deep-space applications (e.g., Fiber Channel, SFODB, ATM, and Myrinet). Moreover, unlike some space applications such as the Mars Pathfinder case, “COTS-based” means the direct use of commercial parts, components, or subsystems, the term COTS has a special meaning for the X2000 architecture. Since at least one of the prospect X2000 customers, namely, Europa, requires to survive in a high-radiation environment, all the critical electronic components must be fabricated on specialized semiconductor foundries. Therefore, another important selection criterion is the availability of the radiation-hardened components for the COTS interface or an ASIC core design (COTS IP) portable to a radiation-hardened foundry. A rigid evaluation based on these criteria resulted in the selection of IEEE 1394. Similar criteria were given to the engineering bus selection while the low-power and performance requirements were further stressed and deemphasized, respectively. Our tradeoff study concluded that I2C was the best compromise¹. The selection of 1394 and I2C enables the X2000 Program to procure COTS ASIC core designs, which can be integrated into a single chip. It is estimated that this approach will reduce the design effort by 30% when compared with the Cassini ASIC design [12], while the complexity of the ASIC is increased by 400%. Moreover, COTS products required by the IEEE 1394 and I2C implementation, such as bus

¹ Although this engineering bus plays an important role in assisting the 1394 system bus for fault detection and recovery, the detailed discussion is beyond the scope of this paper but can be found in [9].

monitors, prototype boards and device drivers, are highly available, which in turn, leads to further savings.

IEEE 1394 is originally designed for commercial applications such as multimedia and portable phones. The current version of IEEE 1394 can support data rates of 100 Mbps, 200 Mbps, and 400 Mbps for the *cable implementation* that is based on a tree topology, and 50 Mbps and 100 Mbps for the *backplane implementation* that is a multi-drop bus. Indeed JPL designers are more familiar with the backplane implementation because of its resemblance to the MIL-STD-1553 bus that was used in the Cassini Project [12]. Nonetheless, we selected the cable implementation due to its high speed and extensive commercial support which enable us to maximize the benefits from using COTS. Accordingly, unless it is explicitly stated, all discussions on IEEE 1394 in this paper refer to the cable implementation.

Although there are various types of tree structure that satisfy the topology requirement of IEEE 1394, it is preferred to have a “regular topology” for space applications. By “regular topology,” we mean a structure that is topologically simple and can be easily maintained as nodes are added to or deleted from the system such that testing and integration can be accomplished efficiently at low cost. Therefore, the stack-tree topology depicted in Figure 1 is proposed, where a node is either a flight computer or a device. There are three physical layer ports in each node. For each branch node, two or more of these ports are connected to other nodes, while a leaf node has only one connected port. Figure 2 depicts the baseline X2000 First Delivery avionics architecture where a stack-tree topology based 1394 dual bus (see Section 3) is shown.




Figure 1: Bus Network based on Stack-Tree Topology




Figure 2: Baseline X2000 First Delivery Architecture

Like any tree structures, the stack-tree topology shown in Figure 1 has a potentially serious drawback. Namely, a tree structure by itself is not fault tolerant as any single node or link failure will result in tree partitioning such that fault-tolerant routing will not be possible. What makes the design more difficult is that to duplicate and cross-strap nodes for bus network fault tolerance purpose are not allowed due to the constraints on power and mass/volume. Although various schemes of fault-tolerant bus network have been proposed in research literatures (see [13, 14], for example), the

restrictions from 1394 and from our application prevent us from utilizing those schemes since the majority of them involve either loops or spare nodes.

The 1394 standard has some error detection provisions such as CRC, they are nonetheless inadequate to ensure reliability for long-life missions. On the other hand, IEEE 1394a [11] provides an employable feature called “port-disable,” which allows us to implement a 1394 compliant reconfigurable bus architecture, though this feature is not intentionally designed for fault tolerance. In the following section, we describe and analyze, the stack-tree topology and its variants based on which we design an IEEE 1394 compliant fault-tolerant bus architecture.

3 Stack-Tree based Bus Architecture

3.1 Concepts

In the interest of bridging the terminology between network topology and the X2000 MCM-stack packaging technology [15], we call the proposed topology “stack-tree topology.”

Definition 1 A stack tree is a tree where each branch node is connected to at most three other nodes among which at most two are branch nodes.

For example, the trees in Figures 3(a), (c) and (d) are stack trees while that in Figure 3(b) is not (as the right node at the first level below the root is connected to three branch nodes).




Figure 3: Trees

Definition 2 A complete stack tree is a stack tree where each branch node is connected to at least one leaf node.

Figure 3(c) depicts a complete stack tree (CST) with n branch nodes. We call this topology *simplex complete stack tree* which is denoted as CST_S . Note that the nodes are labeled such that the branch nodes have the ID numbers from 1 to n , while the leaf nodes have the ID numbers from $n+1$ to $2n$. This labeling scheme will be used in the remainder of the paper. Further, we use n , the number of branch nodes in a CST, to denote the *size* of the tree. Note also that the trees in Figures 3(c) and 3(d) are both CST_S . Based on the CST in Figure 3(c), we can define *CST mirror-image* as follows.

Definition 3 The mirror-image of a complete stack tree is a tree obtained by (1) removing the edges connecting the branch nodes

with the ID numbers i and j which satisfy the relation $|i - j| = 1$; (2) adding edges to connect the leaf nodes with the ID numbers k and l which satisfy the relation $|k - l| = 1$.

Clearly, the CST shown in Figure 3(d) is a mirror image of that in Figure 3(c). It is worth to note that, if we connect $2n$ nodes with a CST-based network and its mirror image, then the two networks will not have any branch nodes in common.

3.2 Applications

3.2.1 The CST_D Scheme

The performance of the X2000 spaceborne systems must be scalable and gracefully degradable. Accordingly, our objective is to develop a fault-tolerant bus network architecture that will allow all the surviving nodes in the bus network to remain connected in the presence of node failures, without requiring spare nodes. The fact that a CST and its mirror image do not have branch nodes in common implies that losing a branch node in one tree will not partition its mirror image. Accordingly, a dual bus scheme comprising a CST and its mirror image, referred to as *CST dual scheme* (denoted as CST_D), as shown in Figure 4(a), will be effective in tolerating single or multiple node failures given that 1) the failed nodes are of the same type (all branch or all leaf) with respect to one of the complete stack trees (see Figure 4(b)), or 2) the failed nodes involve both branch and leaf nodes but they form a cluster at either end (or both ends) of a CST, which will not affect the connectivity of the remainder of the tree (see Figure 4(c)).

We use *terminal clustered branch-leaf failures* to refer to the second failure pattern. Thus, for the cases which involve only the above failure patterns, all the surviving nodes will remain connected (no network partitioning). On the other hand, if a branch node and a leaf node in a CST_D based network fail in a form other than terminal clustered branch-leaf failure (see Figure 4(d)), both the primary and mirror image will be partitioned.

3.2.2 The CST_R Scheme

With the motivation of building a robust bus network architecture capable of tolerating more node failures (in terms of number and pattern), we exploit a unique feature of IEEE 1394, namely, the *port-disable* capability [11]. This feature enables the physical connections between the physical layer of a node and the serial bus cable to become “invisible” from the view point of the remainder of the bus network. The implication is the following:

- 1) By using disabled ports, backup connections between nodes can be added without forming loops (recall that loops are prohibited by IEEE 1394). By “backup connection,” we mean a serial bus cable that connects (via disabled ports) two nodes which are not expected to have a direct connection in the original network configuration (differing from connection replication); and
- 2) Upon fault detection, by disabling physical ports, a failed node will be allowed to be isolated from the rest of the bus network, and necessary backup link(s) can be activated (by enabling the corresponding ports) to repair the partitioned network such that messages can be routed in a reconfigured network, bypassing the failed node.




Figure 4: CST-Based Dual Bus Network (CST_D)

Consider a bus network based on the CST_S topology with n branch nodes (size n), as shown in Figure 5(a). If we add a backup link between any two leaf nodes labeled i and j which satisfy the relation $|(i \bmod n) - (j \bmod n)| = 1$, and also add a backup link to connect branch nodes 1 and n , then we get a topology as shown in Figure 5(b) (an instantiation of the topology in which $n = 6$). Because the added connections (dashed edges) are of inactive nature, the bus network remains free of loop and thus complies with the IEEE 1394 tree topology criterion. Figure 5(c) illustrates the bus network from a 3-dimensional perspective, which enables us to visualize the network as a ring. Accordingly, we denote this bus network configuration as CST_R. To aid the description of failure mechanisms of a CST_R based bus network, we introduce the following terminology:

Definition 4 A failed branch node i and a failed leaf node j in a CST_R based network of size n will form a cut-type failure if $|(j \bmod n) - (i \bmod n)| \leq 1$.




Figure 5: Applying Disabled Backup Links

Figure 6 illustrates the concepts of cut-type and non cut-type failures. Specifically, the failure comprised by nodes 2 and 9 in Figure 6(a), and that by nodes 5 and 11 in Figure 6(b) are cut-type failures. On the other hand, the node failures shown in Figures 6(c) and 6(d) are non cut-type failures. Further, we use the term *clustered failure* to refer to the failure of a group of nodes which are adjacent to each other. Figures 7(a) and 7(b) illustrate the scenarios of clustered and non-clustered multiple cut-type failures, respectively. Clearly, while the non-clustered cut-type failures shown in 7(b) leads to bus network partitioning (i.e., the traffic across either of the “cuts” are disabled), the clustered cut-type failures shown in 7(a) does not even if node 6 also fails (i.e., the traffic across the clustered “cuts” can be re-routed through the enabled backup links $\{1,6\}$ or $\{7,12\}$), although both scenarios involve multiple cut-type failures. The above discussion shows the necessary and

sufficient condition for partitioning a CST_R based bus network:

Remark. A bus network based on the CST_R topology will be partitioned if and only if there exist multiple cut-type failures which do not constitute a single cluster.




Figure 6: Cut-Type and Non Cut-Type Failure

Figure 8 illustrates how partitioning happens in a network that originally has a CST_R structure: The first cut-type failure (single or clustered) will break the ring structure (Figure 8 (a)) so that the remainder of the network becomes a CST_S based structure with backup links (Figure 8 (b)); whereas the second cut-type failure (single or clustered) will break the CST_S based structure, resulting in network partitioning (communication between any two nodes separated by the “cut” becomes impossible, see Figure 8(c)).




Figure 7: Clustered vs. Non-Clustered Multiple Cut-Type Failures

Figure 9 depicts the simplified X2000 architecture in which the CST_R based bus network described above is implemented. In the figure, the solid and dashed thick lines marked “1394 Bus” represent the active and backup links, respectively. During normal operation, the active connections are driven by enabled ports while the ports of backup connections are disabled to avoid loops. The thin lines marked “I2C Bus” correspond to the interface for fault detection, isolation and reconfiguration. The I2C bus is a very simple low-speed multi-drop bus and used only for protecting the




Figure 8: CST_R Based Bus Network Partitioning

1394 bus. Hence this engineering bus has very low utilization and power consumption. For additional protection, a redundant bus (consisting of the 1394 and I2C buses) which is a mirror image of the configuration shown in Figure 9 is proposed by our design [9]. For clarity of illustration, the connections of the redundant bus are not shown in the figure.




Figure 9: CST_R-Based Fault-Tolerant Bus Network

4 Bus Network Reliability Evaluation

4.1 Definition and Notation

In accordance with the objective of the fault-tolerant bus architecture design stated in Section 3.2, we define *bus network reliability* as the probability that, through a mission duration t , the network remains in a state in which all the surviving nodes are connected.

Indeed the causes of a node failure encompass physical layer failure, link layer failure and CPU failure. Moreover, while redundant links (serial bus cables) are permitted in the X2000 architecture, to duplicate and cross-strap nodes for bus network fault tolerance purpose are not allowed due to the power and mass/volume constraints. As a result, the likelihood of node failure is significantly greater than that of link failure. Therefore, in the reliability assessment that follows, we concern only node failure. On the other hand, when a node fails, we assume that there is a possibility that the faulty node may go undetected or the corresponding network reconfiguration process (including port disabling/enabling, etc.) may unexpectedly crash the system. We call the complement of the probability of such an event “coverage.” Before proceeding to derive solutions of reliability measures, we define the following notation:

R_S^{CST} Reliability of a CST_S based bus network.

R_D^{CST} Reliability of a CST_D based bus network.
 R_R^{CST} Reliability of a CST_R based bus network.
 λ Poisson failure rate of a node.
 t Mission duration.
 c Conditional probability (coverage) that a failed node is detected and the corresponding reconfiguration process succeeds given that such a node failure occurs.

4.2 Solutions for R_S^{CST} and R_D^{CST}

We begin with analyzing the CST_S and CST_D based bus network schemes. As explained in Section 3.2, terminal clustered branch-leaf failures in a CST will not affect the connectivity of the remainder of the tree. Thus we can retrieve a “remainder” from the original CST by eliminating the portion(s) comprised by the terminal clustered branch-leaf failures. Consequently, according to the definition of bus network reliability defined in Section 4.1, the reliability calculation for a remainder leads to the solutions of R_S^{CST} and R_D^{CST} . Specifically, we evaluate the reliability of a remainder by conditioning it on its size k (the number of its branch nodes), or equivalently speaking, by conditioning it on the event that the terminal clustered branch-leaf failures involve $(n - k)$ branch-leaf node pairs. Note that a remainder of size k in the original CST of size n has $(n - k + 1)$ possible positions (which in turn, determines the positions of the terminal clustered branch-leaf failures). Note also that the reliability of a remainder of size k for a CST_S based network is the probability that all k branch nodes are failure-free and all faulty leaf nodes are detected and reconfigured successfully (if any). Letting this probability be denoted as $U(k)$, we have,

$$U(k) = (1 - q)^k \sum_{j=0}^k \binom{k}{j} (1 - q)^{k-j} (cq)^j.$$

Then, the theorem of total probability leads to the following solution for R_S^{CST} :

$$R_S^{\text{CST}} = \sum_{k=1}^n (n - k + 1) U(k) (cq)^{2(n-k)}, \quad (1)$$

where $q = 1 - e^{-\lambda t}$ is the probability² that a node fails during mission time t , and $(cq)^{2(n-k)}$ is the probability that $2(n - k)$ nodes are involved in the terminal clustered branch-leaf failures and fault detection and reconfiguration for each of the failed nodes are successful.

Likewise, letting the reliability of a size- k remainder in a CST_D based network be denoted as $V(k)$. According to the failure scenario analysis in Section 3.2.1, we have

$$V(k) = 2(1 - q)^k \sum_{j=1}^k \binom{k}{j} (1 - q)^{k-j} (cq)^j + (1 - q)^{2k}.$$

Thus, the measure R_D^{CST} can be expressed as

$$R_D^{\text{CST}} = \sum_{k=1}^n (n - k + 1) V(k) (cq)^{2(n-k)}. \quad (2)$$

²Although we assume that the time to a node failure is exponentially distributed, the models developed in this paper can accommodate other assumptions such as Weibull distribution by formulating q accordingly.

4.3 A Recursive Model for R_R^{CST}

The solution for R_R^{CST} is impossible to be obtained based on the straightforward use of combinatorics methods because 1) the bus network becomes capable to tolerate more node failure patterns, which makes the representation of the conditions under which the system can survive more complex, and 2) the ring-like structure makes it difficult to ensure that the scenarios considered in the model are exhaustive and mutually exclusive.

As explained in Section 3.2.2, a bus network architecture with a CST_R topology will be partitioned if and only if there exist multiple cut-type failures which do not constitute a single cluster. In other words, the surviving nodes in a CST_R based bus network will remain connected if there exists *at most one* cut-type failure cluster. In the model construction method described below, we view a single cut-type failure as a special case of cut-type failure cluster (where the size of the cluster is one) and treat a size- n network that is free of cut-type failure as a special case of remainder (where the sizes of the cluster and remainder equal to 0 and n , respectively).

Specifically, we first condition network reliability on the size of a cut-type failure cluster (the number of branch nodes involved in the cluster), then we evaluate the probability that the remainder, which is the portion of the CST_R based network structure excluding the cluster, is free of cut-type failure. The key step toward the evaluation of this probability is the derivation of a set of recursive functions that enumerate the combinations and permutations of failed and surviving nodes in the remainder where cut-type failure is absent. By successively expanding and reducing the sizes of the failure cluster and the remainder, respectively, and employing the recursive functions, we exhaustively enumerate the probabilities that a remainder is cut-type failure free. Accordingly, the measure we seek to evaluate can be expressed as

$$R_R^{\text{CST}} = \sum_{m=1}^{n-1} n(cq)^{2m} F(n-m) + F(n) + (n-1)(G(n) + H(n)), \quad (3)$$

where the index m represents the size of a cluster, $(cq)^{2m}$ is the probability that such a cluster exists and each of the individual failed nodes comprising the cluster is detected and undergoes re-configuration successfully, the coefficient n is the number of possible positions of the cluster in the CST_R based network, and $F(n-m)$ evaluates the probability that the remainder is cut-type failure free given that the size of the failure cluster is m . This probability is solved by a set of recursive functions which “walk through” the remainder backward, ensuring that 1) the distinct scenarios characterized by the number and position of failed nodes are exhaustively enumerated, and 2) the remainder is cut-type failure free. More succinctly,

$$F(k) = F_0(k) + F_1(k) + F_2(k), \quad (4)$$

in which k is the size of the remainder ($k = n - m$). Then in each iteration of the recursion, F_0 , F_1 and F_2 individually traverse from the end of a size- i “sub-remainder” where

- 1) both the nodes i and $(i+n)$ are surviving nodes,
- 2) i is a surviving node and $(i+n)$ is a failed node, and
- 3) i is a failed node and $(i+n)$ is a surviving node,

respectively. For each step of the traversal, the recursive functions are derived in a way such that no cut-type failure would be formed within the resulting node-pair sequence. More precisely,

$$\begin{aligned} F_0(i) &= [F_0(i-1) + F_1(i-1) + F_2(i-1)] (1-q)^2(5) \\ F_0(1) &= (1-q)^2 \end{aligned}$$

$$\begin{aligned} F_1(i) &= [F_0(i-1) + F_1(i-1)] (1-q)cq \\ F_1(1) &= (1-q)cq \end{aligned} \quad (6)$$

$$\begin{aligned} F_2(i) &= [F_0(i-1) + F_2(i-1)] cq(1-q) \\ F_2(1) &= cq(1-q) \end{aligned} \quad (7)$$

For example, if nodes i and $(i+n)$ in the remainder are a failed and a surviving node, respectively, then the status of the nodes $(i-1)$ and $(i+n-1)$ must be 1) both surviving, or 2) failed and surviving, respectively (otherwise a cut-type failure will be formed). This particular sequencing rule is implemented by the recursive function F_2 (Equation (7)).

To aid further explanation of the model, we introduce the term *cut-type failure node pair*, abbreviated as *CFP*, which refers to the node pair that forms a cut-type failure. Per Definition 4 in Section 3.2.2, a branch node i in a CST_R based network of size n potentially could be involved in three differing CFPs, as shown in Figure 10. We call the CFPs $\{i, i+n-1\}$, $\{i, i+n\}$ and $\{i, i+n+1\}$ forward CFP, vertical CFP and backward CFP, respectively.




Figure 10: Cut-Type Failure Node Pair

It is worth to note that, due to different patterns of the junctions between the cut-type failure cluster and the remainder, the first term of Equation (3) takes into account not only for the cut-type failure clusters that are constituted by vertical CFPs but also the cut-type failure clusters that can be viewed as the clusters formed by forward and backward CFPs, as illustrated in Figures 11(a) and 11(b), respectively.

For the special case in which $m = 0$, the cut-type failure cluster becomes degenerate while the remainder spans the entire CST_R based bus network. The corresponding scenarios are enumerated by the rest terms in Equation (3), namely, $F(n)$ and $(n-1)(G(n) + H(n))$. Although no coefficient is attached to $F(n)$, different starting positions of this special case “remainder” in the CST_R based bus network are implicitly taken into account by the recursive functions. In other words, as Equations (5), (6) and (7) together exhaustively enumerate the combinations and permutations of failed and surviving nodes such that no cut-type failure will be formed within the remainder, different positions of a “remainder” in the CST_R structure (where $m = 0$) are indeed “inherently” considered by $F(n)$. For example, the “remainder” in Figure 12(b) can be viewed as a result of shifting the starting position of that in Figure 12(a) toward right by one node position — both cases (where $m = 0$) are enumerated by $F(n)$.




Figure 11: Failure Cluster and Remainder




Figure 12: "Positions" of a Remainder of Size n

The only exception is the scenario in which $m = 0$ and a single cut-type failure (a forward or backward CFP) is formed at the "junction" where the two ends of the "remainder" merge, as illustrated in Figure 11(c) (recall that a single CFP by itself will not partition the network). Since the recursive functions are formulated in a way such that the cases where the remainder has an internal CFP are excluded, what missed in the term $F(n)$ but compensated by the term $(n - 1)(G(n) + H(n))$ in Equation (3) are the $(n - 1)$ different positions of this particular CFP. The derivations of $G(n)$ and $H(n)$ are based on two sets of recursive functions which are formulated in a manner such that only the scenarios where the "merge" of the ends of the "remainder" results in a forward CFP and a backward CFP are considered, respectively (while CFP does not exist elsewhere). Due to space limitation, we skip further derivation details.

4.4 Evaluation Results

Applying the models developed in the previous subsections and using *Mathematica*TM, reliability measures for the bus networks based on CST_S , CST_D and CST_R are evaluated with respect to the node failure rate λ , size of bus network n and mission duration t (in hours).

Figure 13 depicts R_S^{CST} , R_D^{CST} and R_R^{CST} as functions of component node failure rate λ . In this evaluation, the size of the CST -based bus networks n is set to 16 (a 32-node network), the fault detection and reconfiguration coverage c is set to 0.9999 (which is conservative as the coverage is defined on a single node basis), and mission duration t is set to 90,000 hours (which implies an *over 10-year* long-life mission). It can be observed from the figure that, while CST_D results in an appreciable amount of improvement from CST_S , CST_R leads to significantly more reliability gain. The quantitative results show that R_R^{CST} will be greater than 0.999997 if node failure rate λ is 10^{-8} or lower. On the other hand, when λ is higher than 10^{-7} , both R_S^{CST} and R_D^{CST} rapidly drop and become unacceptable but R_R^{CST} remains relatively steady.

Figure 14 shows the results of the evaluation for which λ is set to 10^{-7} , t and c remain 90,000 hours and 0.9999, respectively, while n becomes a variable parameter. It is interesting to note that




Figure 13: Bus Network Reliability as a Function of Node Failure Rate ($n = 16$)

R_D^{CST} equals to R_R^{CST} when $n = 2$. This is a reasonable result because for a 4-node network, the node failure patterns that will partition a CST_D based network coincide with the failure patterns that will partition a CST_R based network. It can also be observed that the reliability improvement by R_R^{CST} from R_D^{CST} becomes more significant as the size of the network increases. This is because more routing alternatives that are comprised by active and backup links are available in a larger CST_R based network.

Figure 15 illustrates the evaluation results of a study for which λ and n are set to 10^{-7} and 16, respectively, and c remains 0.9999, while mission duration t becomes a variable parameter. Apparently, both R_S^{CST} and R_D^{CST} become unacceptable for long-life missions. On the other hand, R_R^{CST} remains very reasonable (i.e., 0.999929) even when $t = 100,000$ (a mission duration about 11.5 years).

5 Conclusion

We have presented qualitative and quantitative analyses of a COTS-based fault-tolerant bus network architecture. To implement COTS-based fault tolerance is becoming a major challenge for us today when cost concern has led to increased use of COTS products for critical applications. On the other hand, vendors remain reluctant to incorporate fault tolerance features into COTS products because doing so is likely to increase development and production costs and thus weaken the market competitiveness of their products. Therefore, to cope with the current state of COTS




Figure 14: Bus Network Reliability as a Function of Network Size ($\lambda = 10^{-7}$)




Figure 15: Bus Network Reliability as a Function of Mission Duration ($\lambda = 10^{-7}$)

is crucial for us. Our analyses demonstrate that rigid assessment and innovative utilization of pertinent standard features of a COTS product could enable us to circumvent its shortcomings and facilitate us to implement COTS-based fault-tolerance for critical applications. Further, our effort on COTS-based fault tolerance reported in this paper and the developments of COTS-based highly reliable systems by other organizations suggest to the vendors the following: It could be indeed beneficial to them to incorporate fault tolerance features into COTS products as implementation options. This potential benefit stems from the fact that the implementation options will permit a COTS product, in a cost-effective manner, to satisfy both critical and non-critical application customers, strengthening the market competitiveness of the product.

From analytic method point of view, the recursive function based model described in this paper enables us to obtain the exact solution of the reliability measure R_R^{CST} that is otherwise very difficult to solve. Indeed, our recursive function based method can be rather easily adapted for the evaluation of the bus architec-

tures which have more constraints on reconfiguration and performance degradation. Aimed at further reliability and performance improvements, we are currently developing an extended model for a comprehensive study of the CST_R based bus architecture, using more sophisticated measures including performability.

References

- [1] L. Alkalai and A. T. Tai, "Long-life deep-space applications," *IEEE Computer*, vol. 31, pp. 37–38, Apr. 1998.
- [2] L. Alkalai and S. N. Chau, "Description of X2000 avionics program," in *Proceedings of the 3rd DARPA Fault-Tolerant Computing Workshop*, (Pasadena, CA), June 1998.
- [3] A. T. Tai, L. Alkalai, and S. N. Chau, "On-board preventive maintenance: A design-oriented analytic study for long-life applications," *Performance Evaluation*, vol. 35, pp. 215–232, June 1999.
- [4] L. Alkalai, "A roadmap for space microelectronics technology into the New Millennium," in *Proceedings of the 35th Space Congress*, (Cocoa Beach, FL), Apr. 1998.
- [5] T. Shanley and D. Anderson, *PCI System Architecture*. Addison Wesley, 1995.
- [6] IEEE 1394, *Standard for a High Performance Serial Bus*. Institute of Electrical and Electronic Engineers, Jan. 1995.
- [7] D. Anderson, *FireWire System Architecture, IEEE 1394*. PC System Architecture Series, MA: Addison Wesley, 1998.
- [8] D. Paret and C. Fenger, *The I2C Bus: From Theory to Practice*. John Wiley, 1997.
- [9] S. N. Chau, L. Alkalai, J. B. Burt, and A. T. Tai, "The design of a fault-tolerant COTS-based bus architecture," in *Proceedings of 1999 Pacific Rim International Symposium on Dependable Computing (PRDC'99)*, (Hong Kong, China), Dec. 1999.
- [10] J. R. Marshall, "Building standards based COTS multiprocessor computer systems for space around a high speed serial bus network," in *Proceedings of the 17th Digital Avionics Systems Conference*, (Bellevue, Washington), Nov. 1998.
- [11] IEEE P1394A, *Standard for a High Performance Serial Bus (Supplement), Draft 2.0*. Institute of Electrical and Electronic Engineers, Mar. 1998.
- [12] J. Donaldson, "Cassini Orbiter Functional Requirements Book: Command and Data Subsystem," JPL Document CAS-4-2006, Jet Propulsion Laboratory, California Institute of Technology, Pasadena, CA, June 1994.
- [13] C. S. Raghavendra, A. Avižienis, and M. D. Ercegovic, "Fault tolerance in binary tree architecture," *IEEE Trans. Computers*, vol. C-33, pp. 568–572, June 1984.
- [14] Y.-R. Leu and S.-Y. Kuo, "A fault-tolerant tree communication scheme for hypercube systems," *IEEE Trans. Computers*, vol. C-45, pp. 641–650, June 1996.
- [15] K. Sasidhar, L. Alkalai, and A. Chatterjee, "Testing NASA's 3D-stack MCM space flight computer," *IEEE Design & Test of Computers*, vol. 15, pp. 44–55, July-September 1998.